

Building Image Classification Models



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

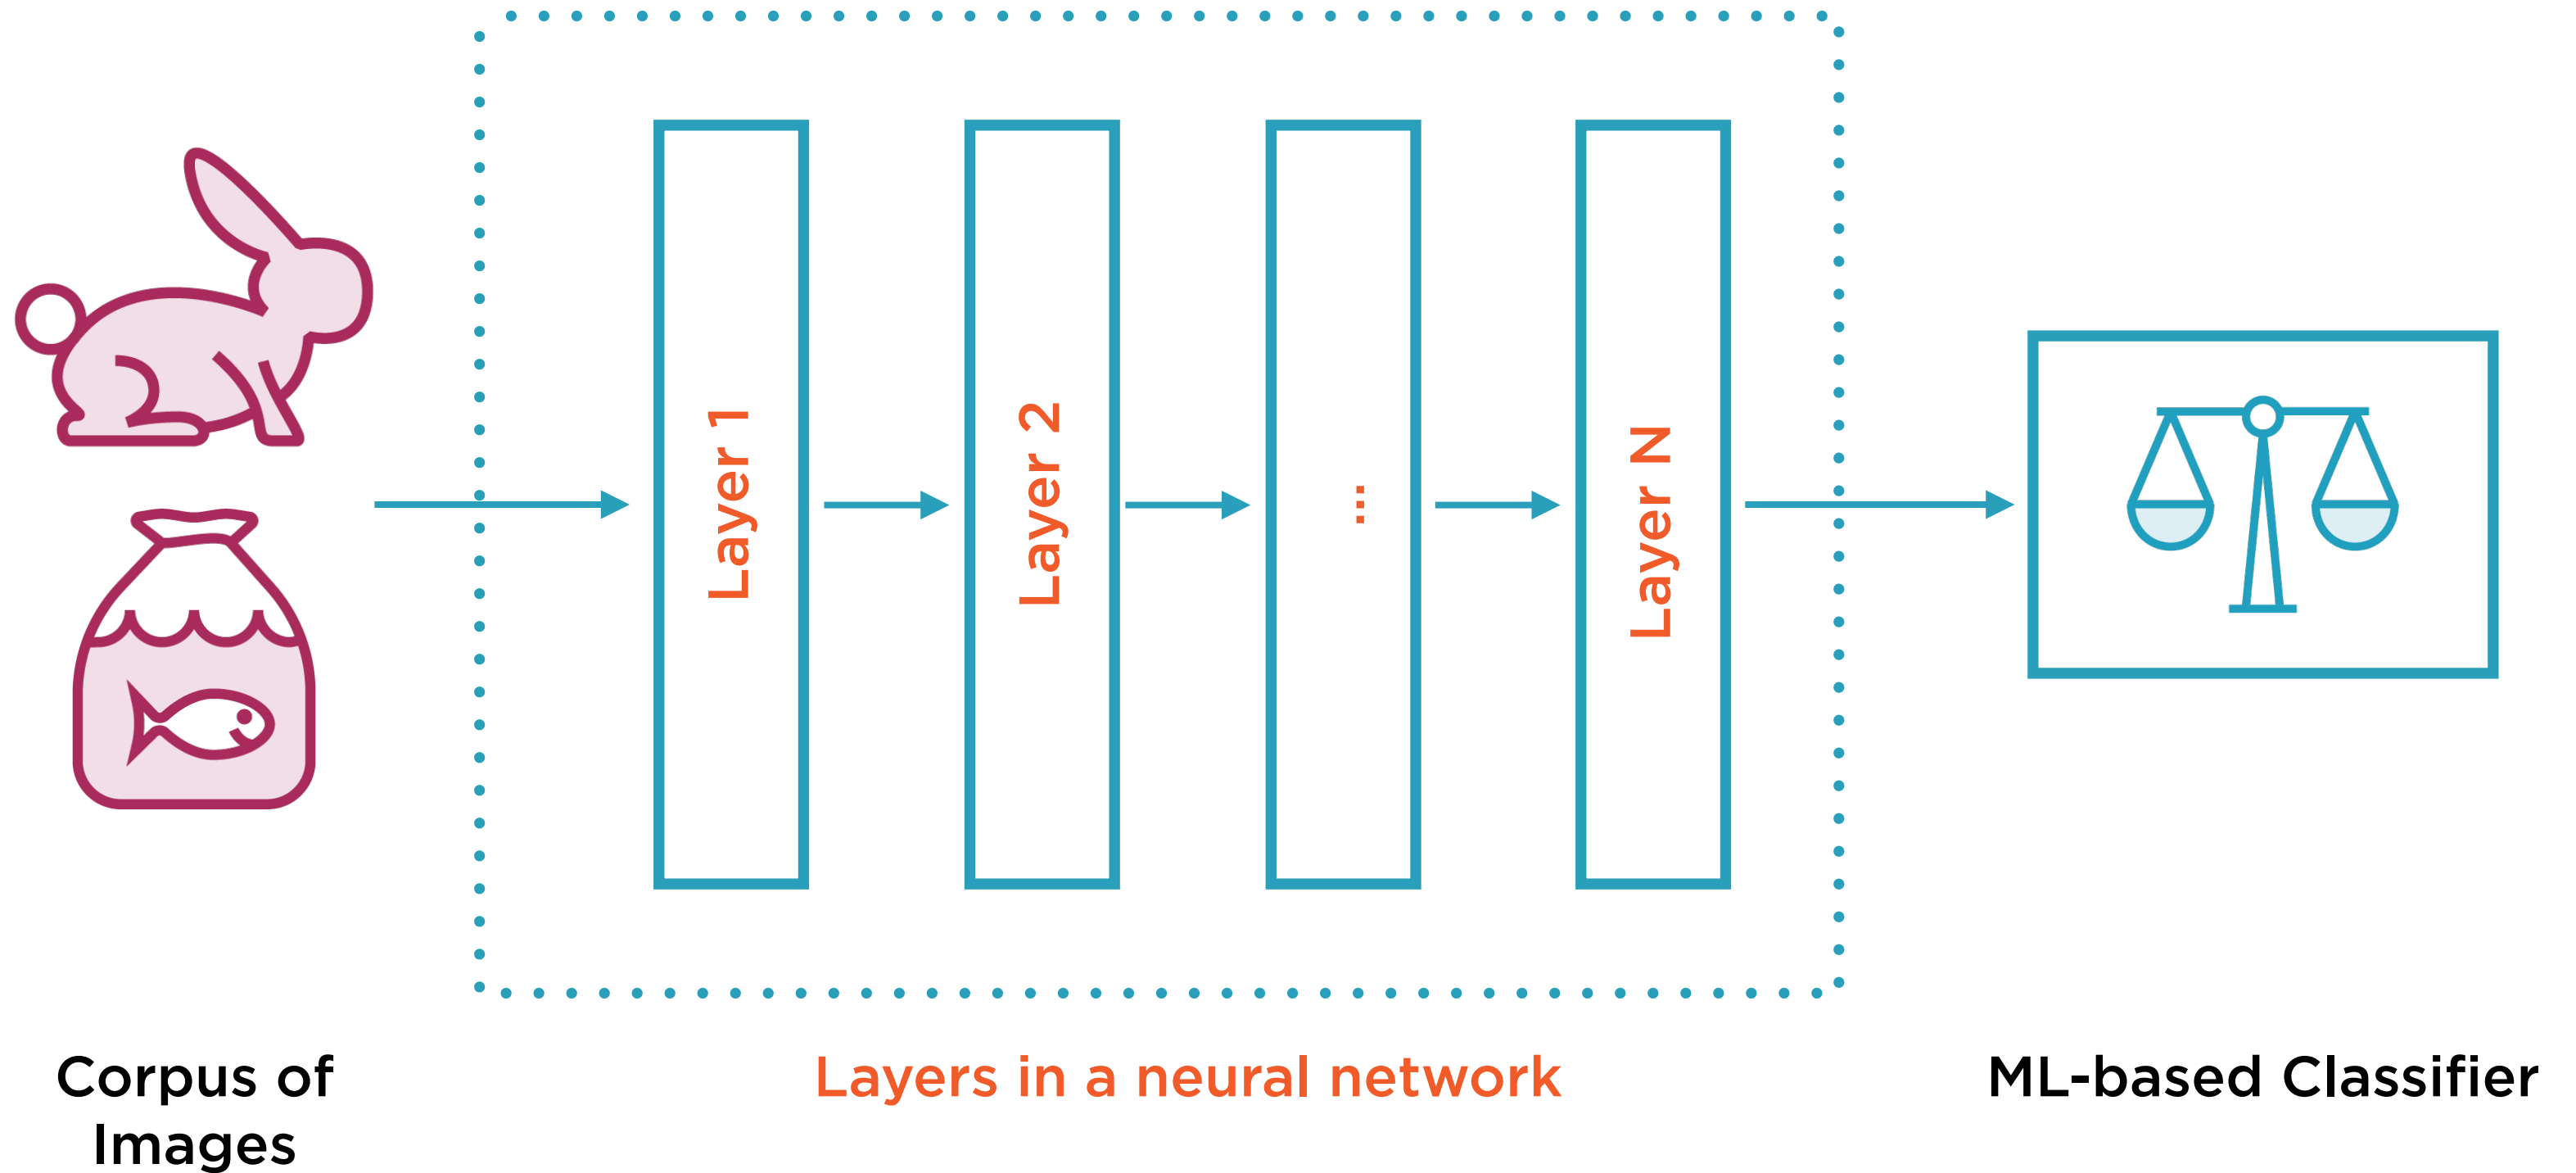
Image classification models

Convolutional layers and pooling layers

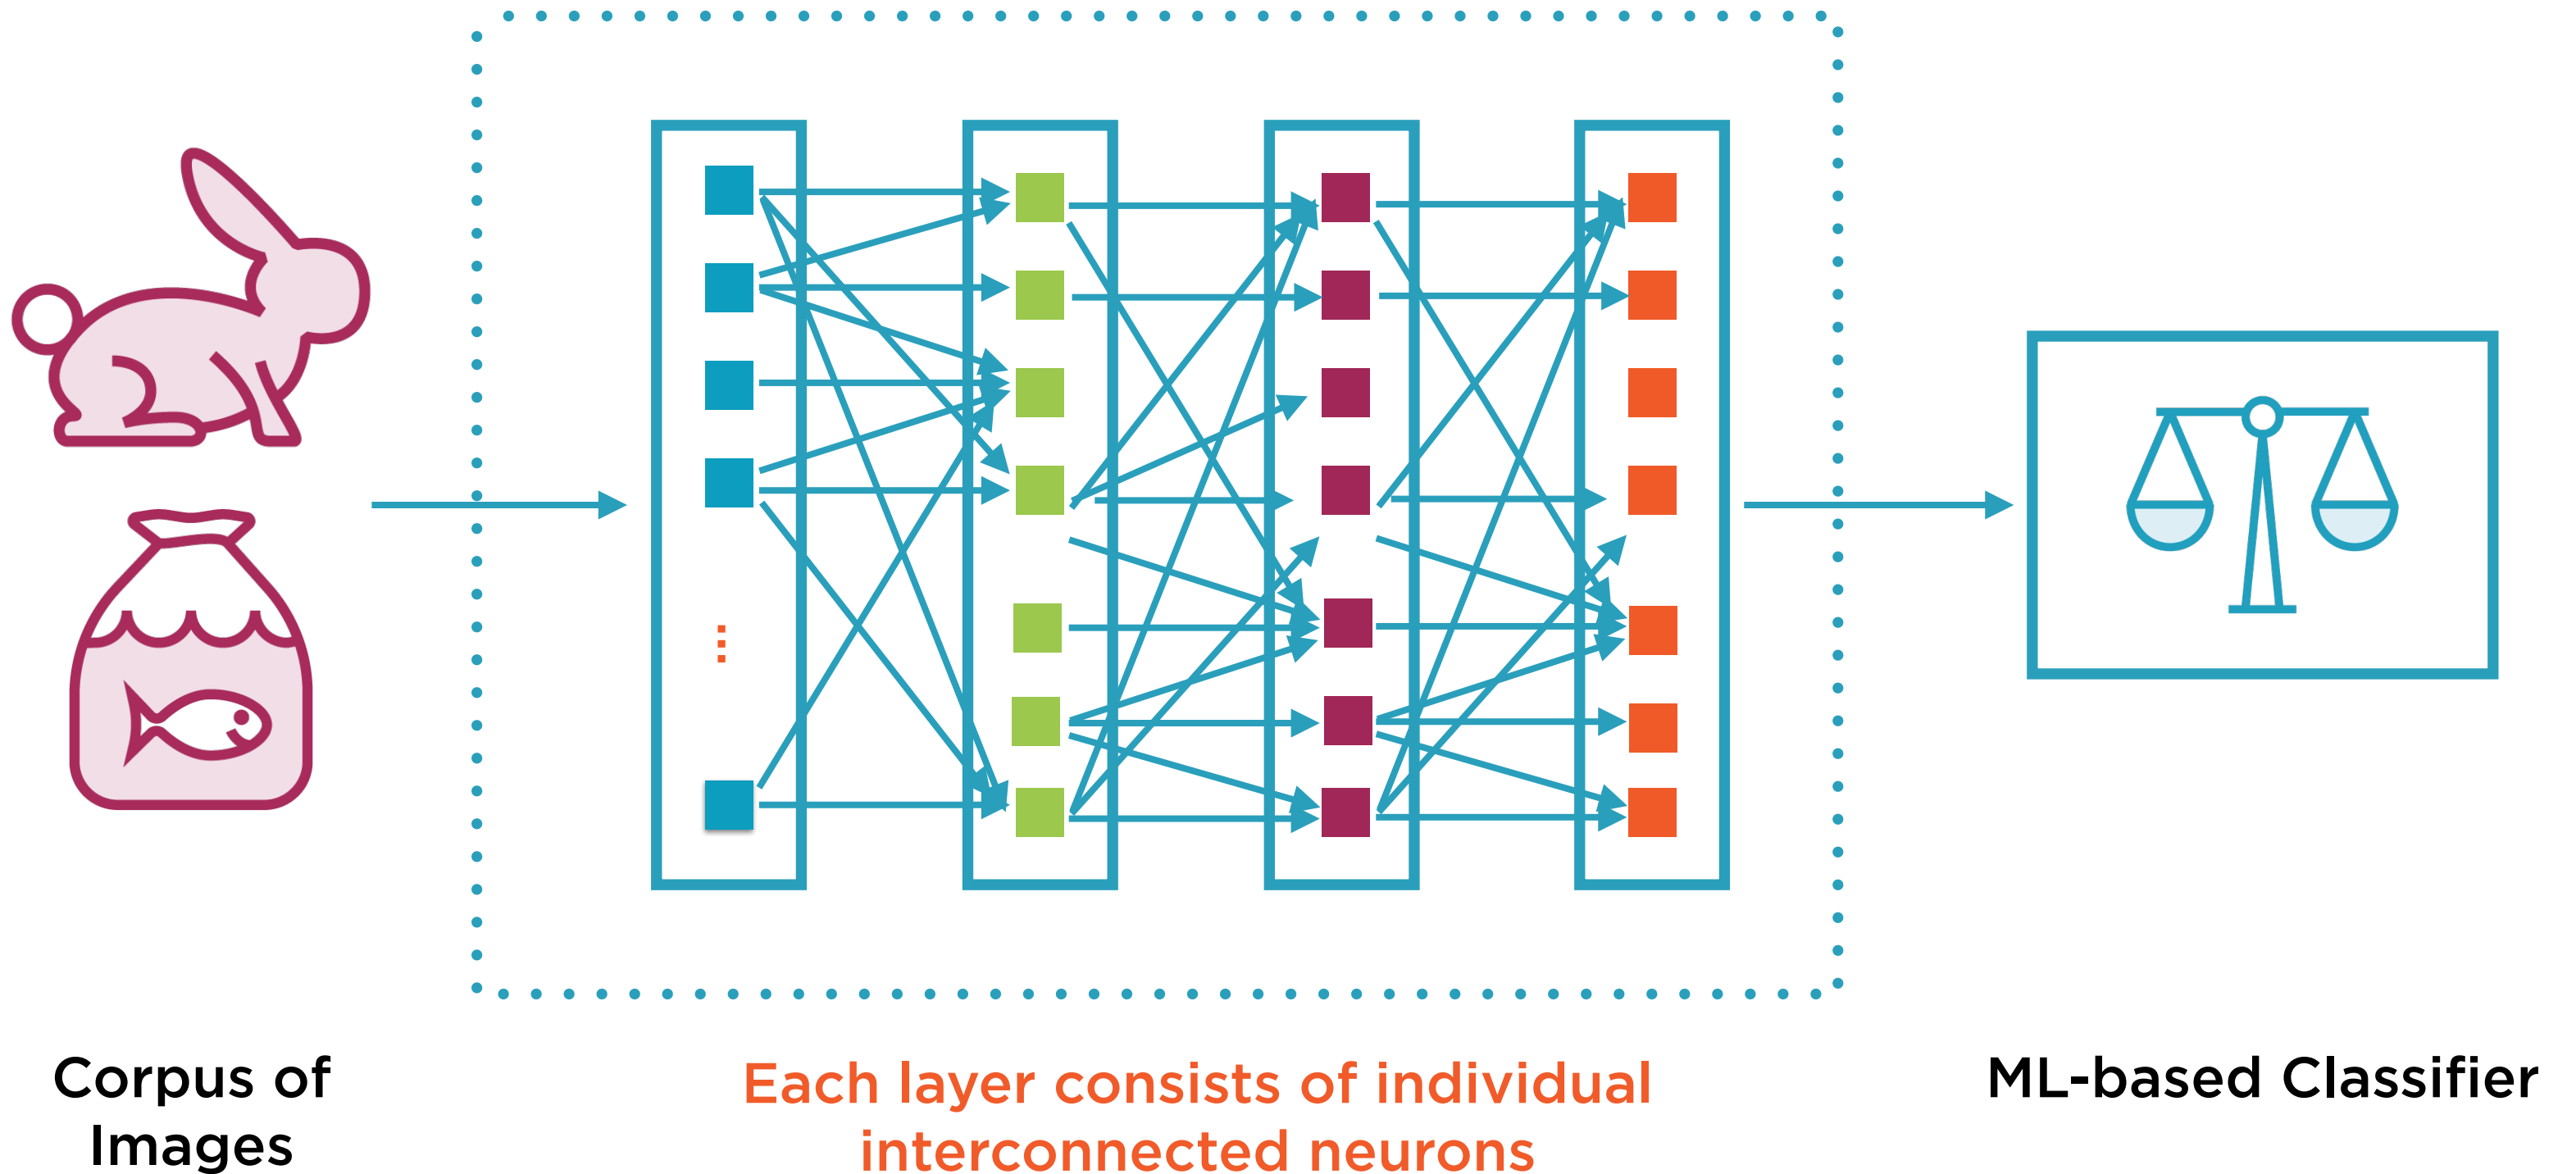
**Convolutional Neural Networks (CNNs)
for image classification**

**Implementing CNNs in tf.keras for
image classification**

Neural Networks for Image Classification



Neural Networks for Image Classification



Parameter Explosion



**Consider a 100 x 100 pixel image
(10,000 pixels)**

If first layer = 10,000 neurons

Interconnections ~ $O(10,000 * 10,000)$

**100 million parameters to train neural
network!**

Parameter Explosion



**Dense, fully connected neural networks
can't cope**

**Also do not provide feature extraction
with location invariance**

**Convolutional neural networks to the
rescue**

Introducing Convolutional Neural Networks

Dense neural networks do
not consider the spatial
aspects of images

Viewing an Image




All neurons in the eye don't see the entire image

Viewing an Image



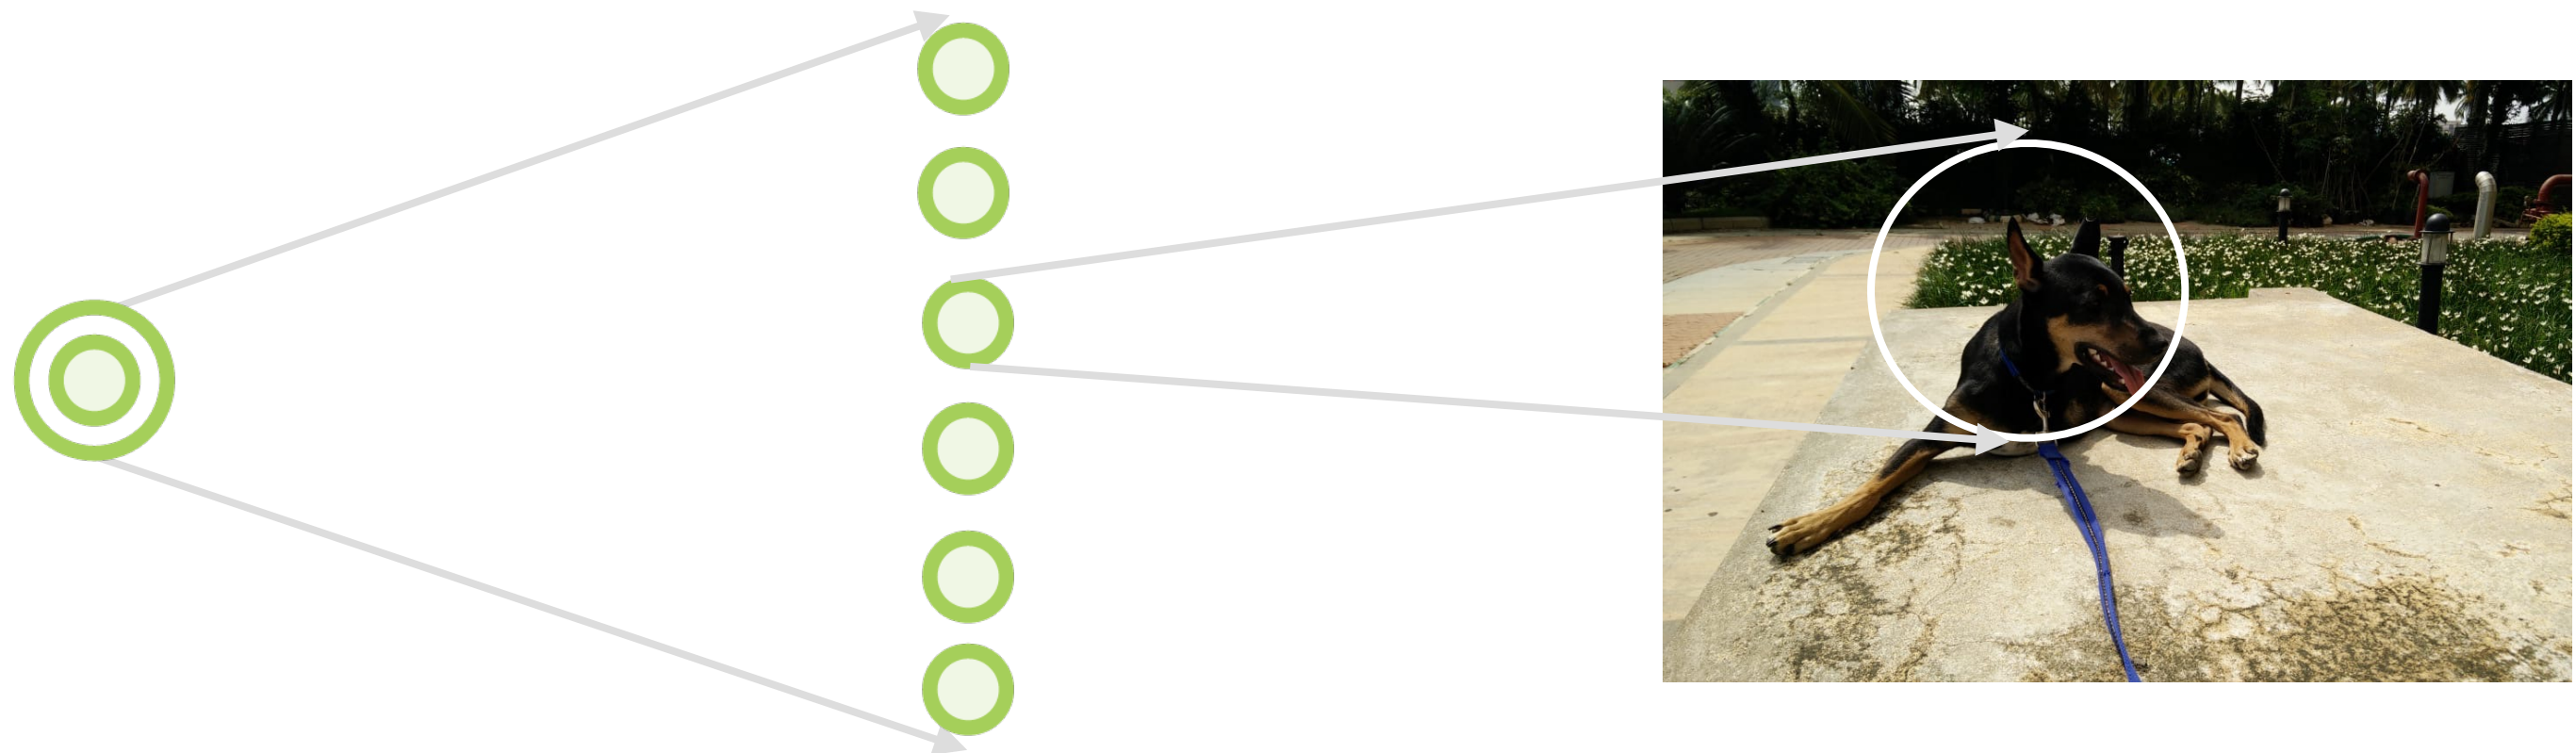
Each neuron has its own local receptive field

Viewing an Image




**It reacts only to visual stimuli located in its
receptive field**

Viewing an Image



Some neurons react to more complex patterns
that are **combinations** of lower level patterns

Neural Networks



Sounds like a classic neural network problem

Convolutional neural networks
consider the **spatial** aspects of
image and **aggregate**
information from local fields

Convolutional Neural Networks



Eye perceives visual stimulus in 2D visual field

“Local receptive field”

Eye sends 2D image to visual cortex

Convolutional Neural Networks



Visual cortex adds depth perception

Individual neurons in cortex focus on small field

Convolutional Neural Networks



CNNs perform spectacularly well at many tasks

Particularly at image recognition


Dramatically fewer parameters than DNN with similar performance

Inspirations for CNNs



Two Dimensions

Data comes in expressed
in 2D



Local Receptive Fields

Neurons focus on narrow
portions

Two Kinds of Layers in CNNs

Convolution

Local receptive field

Pooling

Subsampling of inputs

Convolution

Two Kinds of Layers in CNNs

Convolution

Local receptive field

Pooling

Subsampling of inputs

Convolution

In this context, a sliding window function applied to a matrix

Convolution

In this context, a sliding window function applied to
a matrix



e.g. a matrix of pixels
representing an image

Convolution

In this context, a sliding window **function** applied to a matrix

Often called a kernel or filter




Convolution

In this context, a sliding window function applied to a matrix




Kernel is applied element-wise
in sliding-window fashion

Representing Images as Matrices



= 784 pixels

Representing Images as Matrices




= 36 pixels

Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix




Kernel

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



Kernel

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



x1	x0	x1
x0	x1	x0
x1	x0	x1

4

4

1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution
Result

Convolution

0 _{x1}	x0	0 _{x1}	0	0	0
x0	0.8 _{x1}	x0	0.3	0.6	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



Convolution
Result

Convolution

0_{x1}	$x0$	0_{x1}	0	0	0
$x0$	0.8_{x1}	$x0$	0.3	0.6	0
0.2_{x1}	$x0$	0_{x1}	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution
Result

Convolution

0	0 _{x1}	x0	0 _{x1}	0	0
0.2	x0	0 _{x1}	x0	0.6	0
0.2	0.9 _{x1}	x0	0.3 _{x1}	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution
Result

Convolution

0	0 _{x1}	x0	0 _{x1}	0	0
0.2	x0	0 _{x1}	x0	0.6	0
0.2	0.9 _{x1}	x0	0.3 _{x1}	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution
Result

Convolution

0	0	0 _{x1}	x0	0 _{x1}	0
0.2	0.8	x0	0.3 _{x1}	x0	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution
Result

Convolution

0	0	0 _{x1}	x0	0 _{x1}	0
0.2	0.8	x0	0.3 _{x1}	x0	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution
Result

Convolution

0	0	0	0 _{x1}	x0	0 _{x1}
0.2	0.8	0	x0	0.6 _{x1}	x0
0.2	0.9	0	0.3 _{x1}	x0	0 _{x1}
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution
Result

Convolution

0	0	0	0 _{x1}	x0	0 _{x1}
0.2	0.8	0	x0	0.6 _{x1}	x0
0.2	0.9	0	0.3 _{x1}	x0	0 _{x1}
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution
Result

Convolution

0	0	0	0	0	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.6	0
x0	0.9 _{x1}	x0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution
Result

Convolution

0	0	0	0	0	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.6	0
x0	0.9 _{x1}	x0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8 _{x1}	x0	0.3 _{x1}	0.6	0
0.2	x0	0 _{x1}	x0	0.8	0
0.3	0.8 _{x1}	x0	0.8 _{x1}	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8 _{x1}	x0	0.3 _{x1}	0.6	0
0.2	x0	0 _{x1}	x0	0.8	0
0.3	0.8 _{x1}	x0	0.8 _{x1}	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0 _{x1}	x0	0.6 _{x1}	0
0.2	0.9	x0	0.3 _{x1}	x0	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0 _{x1}	x0	0.6 _{x1}	0
0.2	0.9	x0	0.3 _{x1}	x0	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 _{x1}	x0	0 _{x1}
0.2	0.9	0	x0	0.8 _{x1}	x0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 _{x1}	x0	0 _{x1}
0.2	0.9	0	x0	0.8 _{x1}	x0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.8	0
x0	0.8 _{x1}	x0	0.8	0.9	0
0 _{x1}	x0	0 _{x1}	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.8	0
x0	0.8 _{x1}	x0	0.8	0.9	0
0 _{x1}	x0	0 _{x1}	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9_{x1}	x0	0.3_{x1}	0.8	0
0.3	x0	0.7_{x1}	x0	0.9	0
0	0_{x1}	x0	0.2_{x1}	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9_{x1}	x0	0.3_{x1}	0.8	0
0.3	x0	0.7_{x1}	x0	0.9	0
0	0_{x1}	x0	0.2_{x1}	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	x0	0.8 _{x1}	x0	0
0	0	0 _{x1}	x0	0.8 _{x1}	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	x0	0.8 _{x1}	x0	0
0	0	0 _{x1}	x0	0.8 _{x1}	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3_{x1}	x0	0_{x1}
0.3	0.8	0.7	x0	0.9_{x1}	x0
0	0	0	0.2_{x1}	x0	0_{x1}
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3 _{x1}	x0	0 _{x1}
0.3	0.8	0.7	x0	0.9 _{x1}	x0
0	0	0	0.2 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
x0	0 _{x1}	x0	0.2	0.8	0
0 _{x1}	x0	0 _{x1}	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
x0	0 _{x1}	x0	0.2	0.8	0
0 _{x1}	x0	0 _{x1}	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8_{x1}	x0	0.8_{x1}	0.9	0
0	x0	0_{x1}	x0	0.8	0
0	0_{x1}	x0	0.2_{x1}	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8 _{x1}	x0	0.8 _{x1}	0.9	0
0	x0	0 _{x1}	x0	0.8	0
0	0 _{x1}	x0	0.2 _{x1}	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7_{x1}	x0	0.9_{x1}	0
0	0	x0	0.2_{x1}	x0	0
0	0	0_{x1}	x0	0.2_{x1}	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	x0	0.2 _{x1}	x0	0
0	0	0 _{x1}	x0	0.2 _{x1}	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	x0	0.8 _{x1}	x0
0	0	0	0.2 _{x1}	x0	0 _{x1}

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	x0	0.8 _{x1}	x0
0	0	0	0.2 _{x1}	x0	0 _{x1}

Matrix




1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution
Result

Convolutional Layers

Convolutional Layers



Convolution layers - zoom in on specific bits of input

Extract structure and features in the input image


Successive layers aggregate inputs into higher level features

Pixels >> Lines >> Edges >> Object


Feature Maps



Image




Pixels




Feature
Map

Feature maps are convolutional layers generated by applying a convolutional kernel to the input


Feature Maps




Pixels



Convolutional
Layer




Neurons



Number of neurons
in receptive field =
kernel size

Feature Maps


Neuron i



Pixels

Convolutional
Layer

Kernel Size

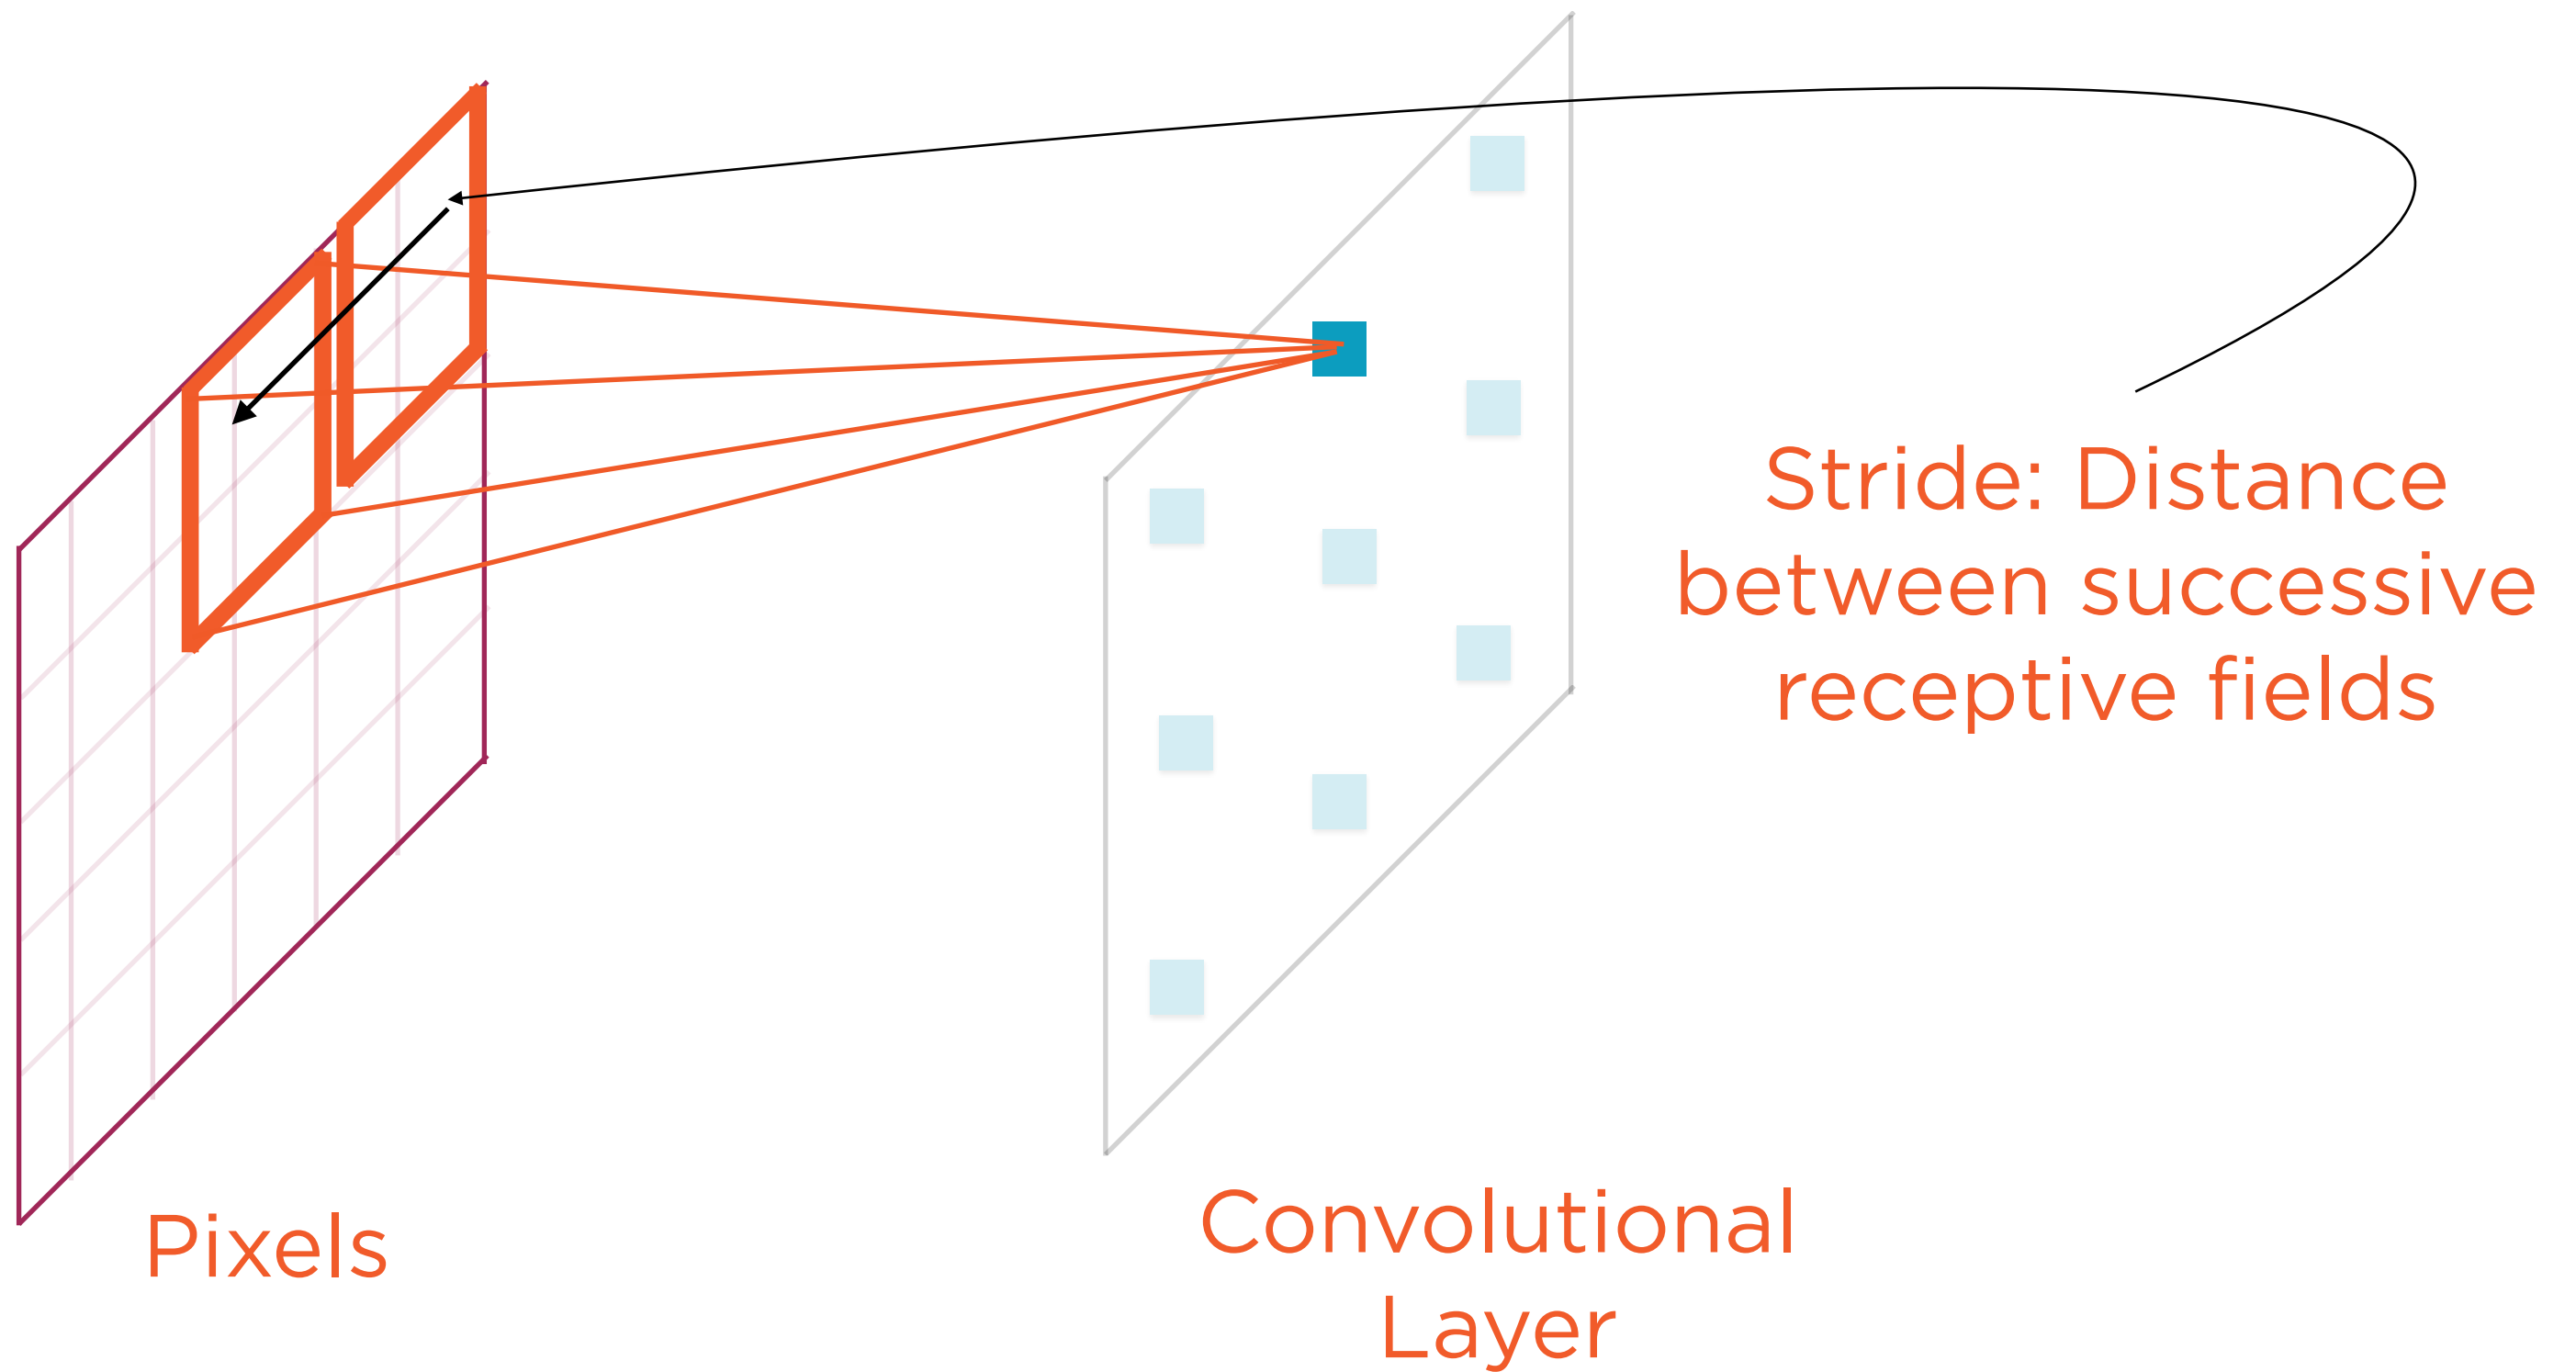


Convolutional kernel size usually expressed in terms of width and height of receptive area

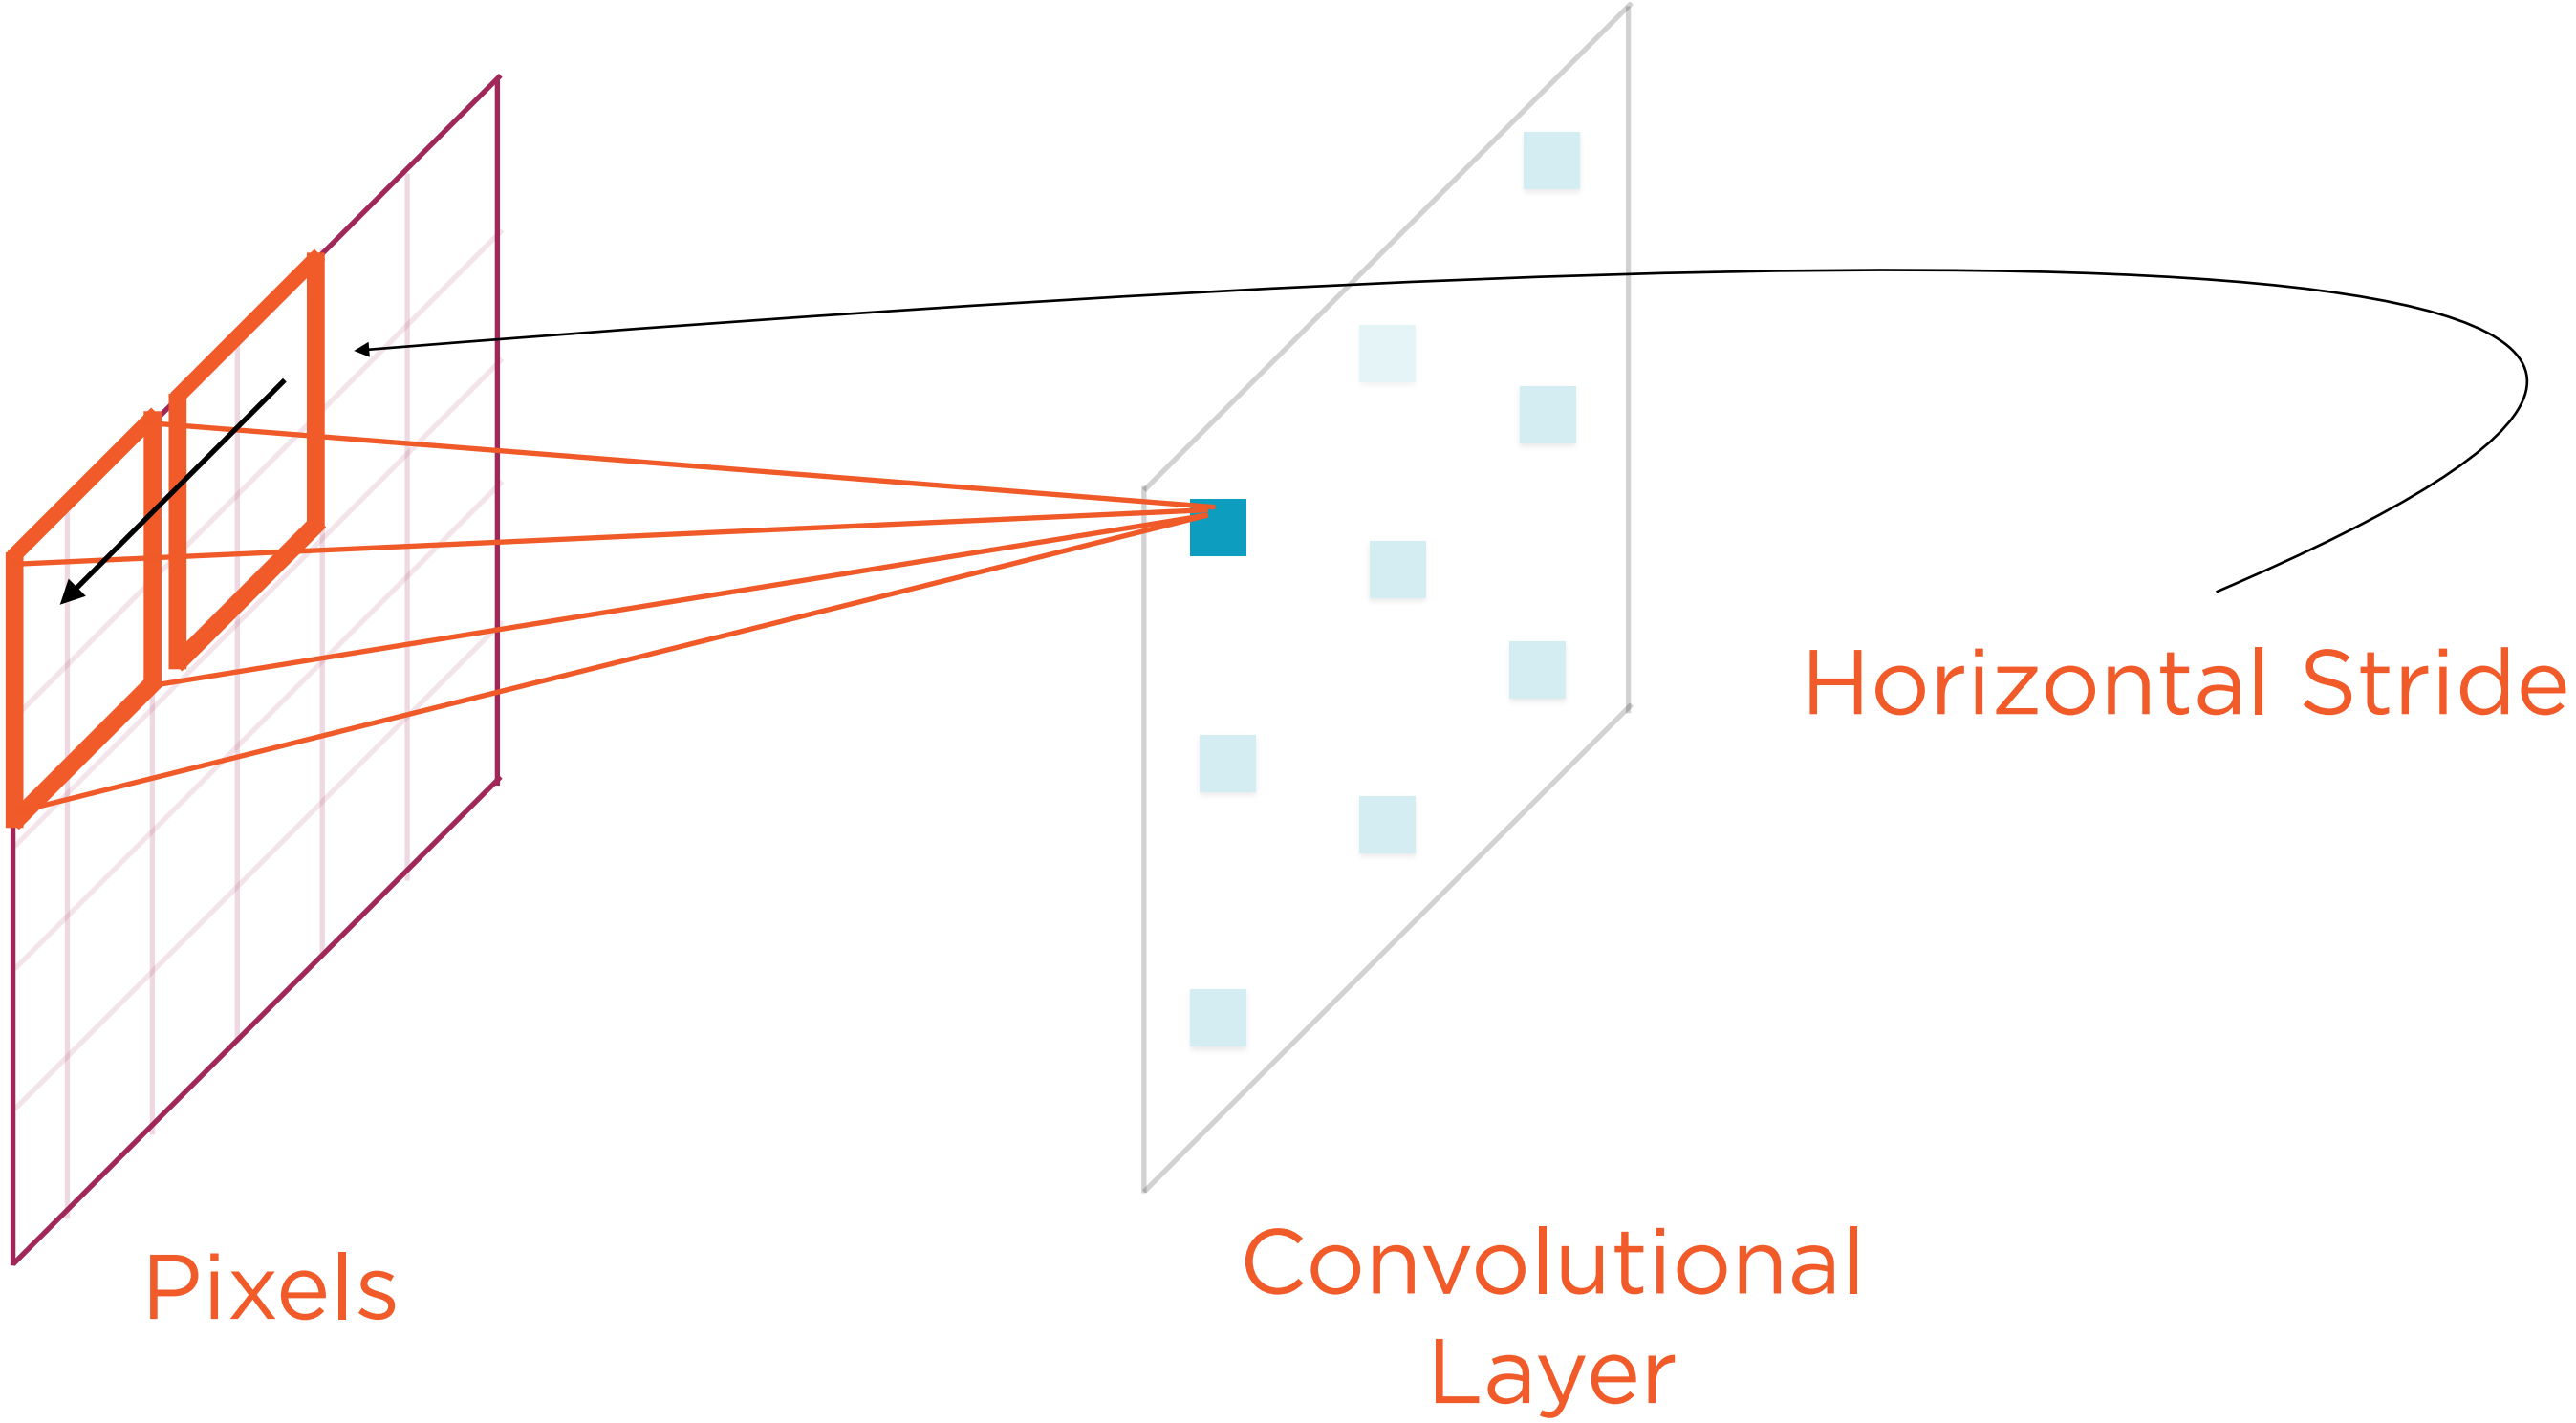
Use **small convolutional kernels, more efficient**

Stacking two 3x3 kernels is preferable to one 9x9 kernel


Feature Maps




Feature Maps




Feature Maps



Feature Maps




Feature Maps




Feature Maps

Zero padding
may be needed
at the edges




Feature Maps



Sparse, not
Dense

Feature Maps




All neurons in a feature map have the same weights and biases

Two big advantages over DNNs

- Dramatically **fewer** parameters to train
- CNN can recognize feature patterns **independent** of location

Feature Maps




The parameters of all neurons in a feature map are collectively called the filter


Why filter?

Because weights **highlight (filter)** specific patterns from the input pixels


CNNs



Feature
Map




Convolutional
Layer



CNN

Convolutional Layer



Each convolutional layer consists of several feature maps of equal sizes

The different feature maps have different parameters

Pooling Layers

Two Kinds of Layers in CNNs


Convolution

Local receptive field

Pooling

Subsampling of inputs


Pooling



Matrix




Max,
2x2 filter,
stride = 2



Pooling Result


Pooling



Matrix




Max,
2x2 filter,
stride = 2



Pooling Result


Pooling



Matrix




Max,
2x2 filter,
stride = 2



Pooling Result


Pooling



Matrix




Max,
2x2 filter,
stride = 2




Pooling Result

Pooling




Matrix

Max,
2x2 filter,
stride = 2



Pooling Result

Pooling Layers




Neurons in a pooling layer have no weights or biases

A pooling neuron simply applies some aggregation function to all inputs

Max, sum, average


Pooling Layers



Why use them?

- Greatly reduce memory usage during training
- Mitigate overfitting (via subsampling)
- Make NN recognize features independent of location (location invariance)

Pooling Layers




Pooling layers typically act on each channel independently

So, usually, output area < input area but

Output depth = Input depth


CNN Architectures

Typical CNN Architecture




Alternating groups of convolutional and pooling layers

Typical CNN Architecture




**Each group of convolutional layers usually
followed by a ReLU layer**

Typical CNN Architecture




The output of each layer is also an image

Typical CNN Architecture




However successive outputs are smaller and smaller (due to pooling layers)

Typical CNN Architecture



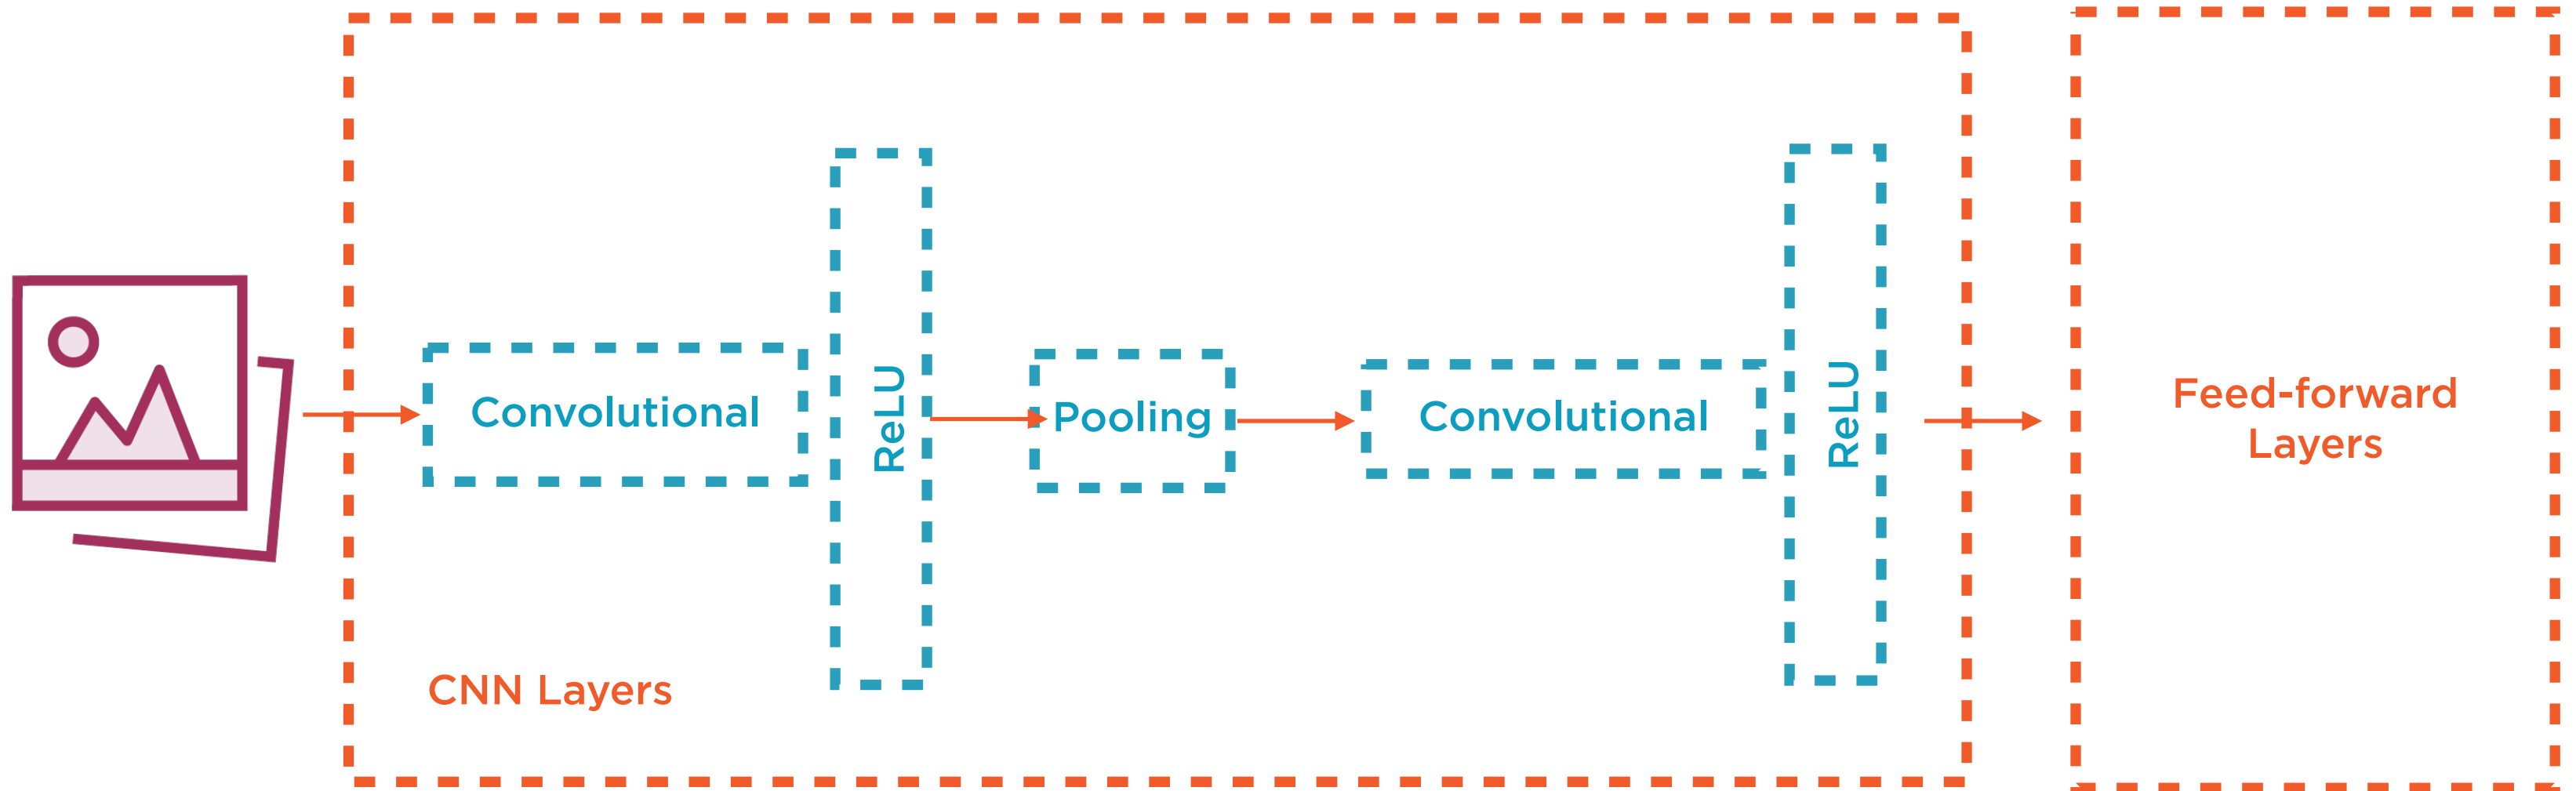
**As well as deeper and deeper (due to feature maps
in the convolutional layers)**

Typical CNN Architecture




This entire set of layers is then fed into a regular, feed-forward neural network

Typical CNN Architecture



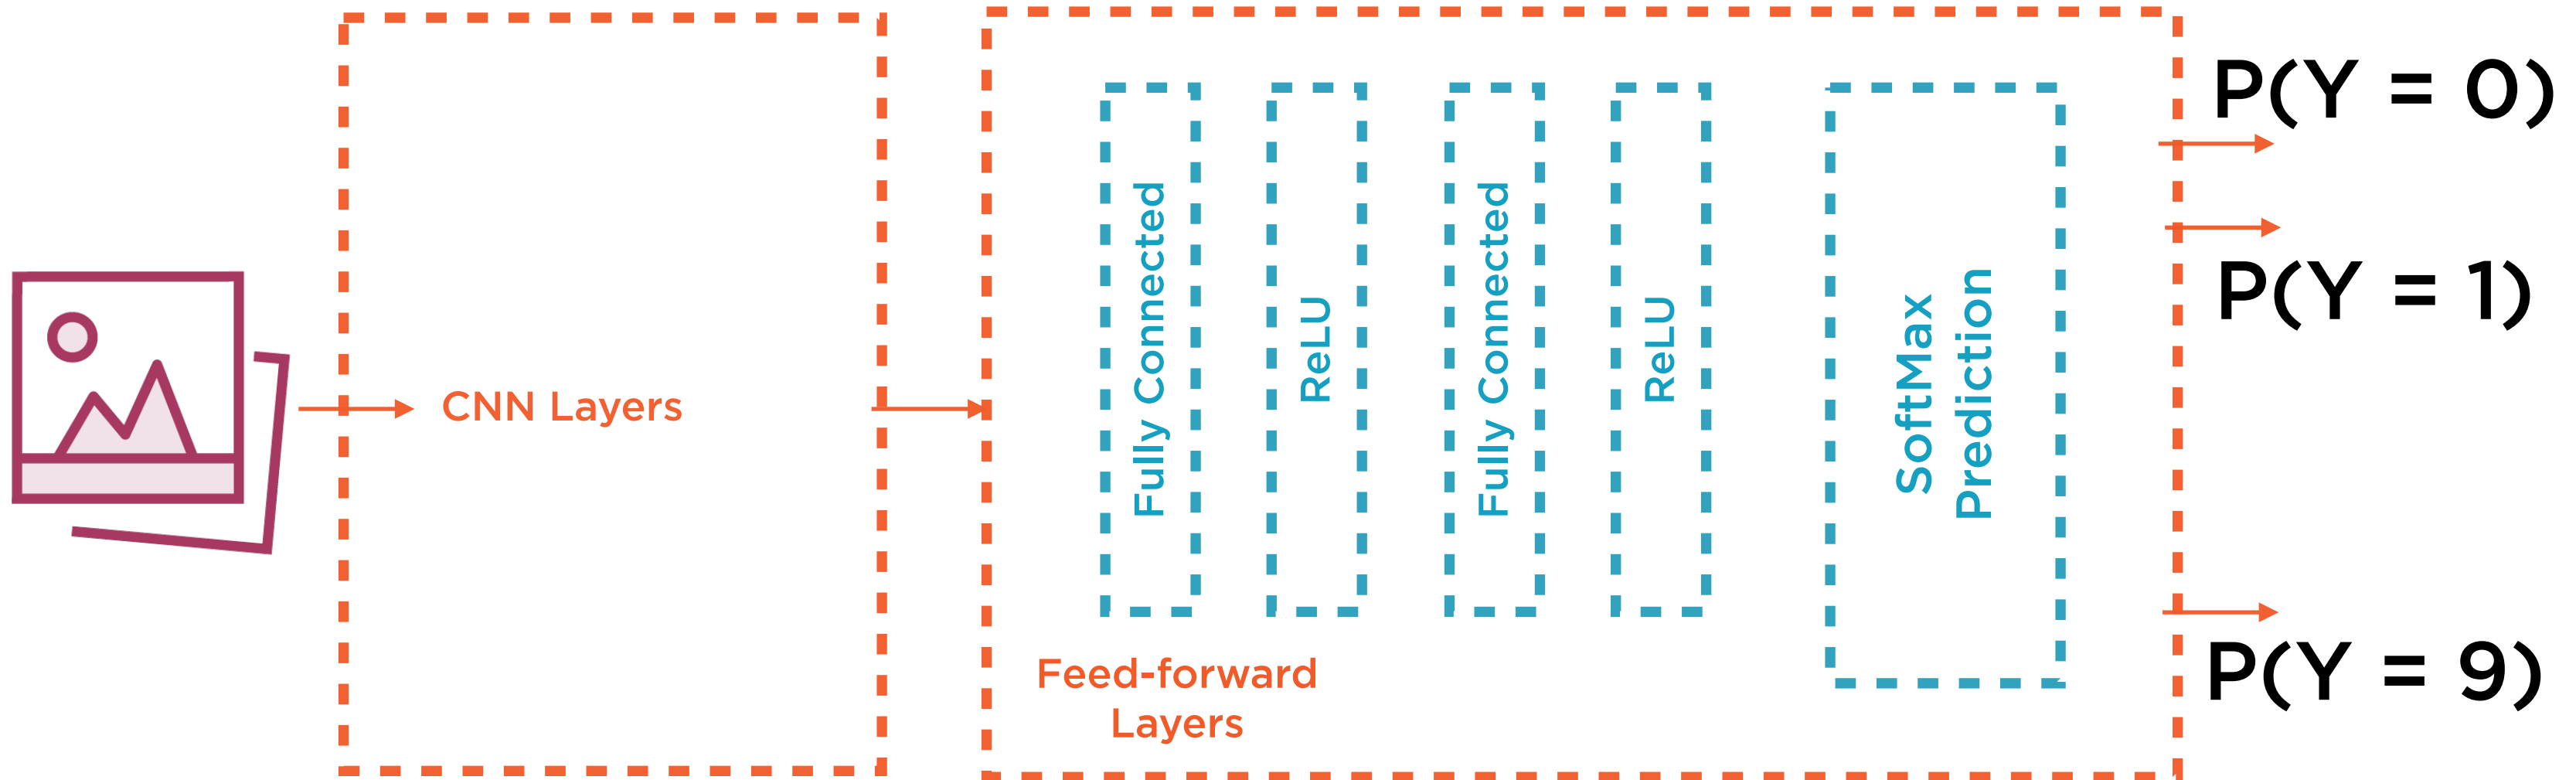
**This entire set of layers is then fed into a regular,
feed-forward neural network**

Typical CNN Architecture




**This feed-forward has a few fully connected layers
with ReLU activation**

Typical CNN Architecture



This is the output layer, emitting probabilities

Typical CNN Architectures



Input is an image

Outputs are probabilities

Demo

**Image classification using a
convolutional neural network**

Summary

Image classification models

Convolutional layers and pooling layers

**Convolutional Neural Networks (CNNs)
for image classification**

**Implementing CNNs in tf.keras for
image classification**

Up Next:

Building Unsupervised Machine
Learning Models
